

1 Microbiology

A laboratory wants to determine if two different methods (A and B) give similar results for quantifying a particular bacterial species in a particular medium. Under each method, the counts form a random sample. Assume that the counts follow a Poisson distribution, since this distribution is a typical model for such data. Let μ_A and μ_B represent the population mean counts for Methods A and B, respectively. Let $\theta = \mu_A \mu_B$.

1.1 Bootstrap estimate of SE

Use the bootstrap to estimate the standard error of $\hat{\theta} = \mu_A - \mu_B$, where μ_A and μ_B are sample means of counts for methods A and B, respectively. Use $B = 2000$ bootstrap samples. Assume the replicates are **unpaired**.

- Let $A = X$ and $Y = B$. I am changing the labeling so as not to confuse with the number of bootstrap samples, B .
- sampling is unpaired
- Poisson distribution is assumed. Hence, this is parametric bootstrap. Recall:

$$\mathcal{P}(X = x \mid \lambda) = \frac{e^{-\lambda} \lambda^x}{x!}; \quad x = 0, 1, \dots; \quad 0 \leq \lambda < \infty \quad (1.1)$$

- According to this reference. A sufficient statistic for the parameter of the Poisson distribution, λ is $\sum_{i=1}^n x_i$. Since it is sufficient to know the average if we know the sum (given the sample), then we can also take the mean as another sufficient statistic.

Algorithm:

Step 1: Calculate a sufficient statistic for λ_X (method A) and λ_Y (method B).

$$\begin{aligned} \hat{\lambda}_X &= \frac{\sum_{i=1}^n x_i}{n} \\ \hat{\lambda}_Y &= \frac{\sum_{i=1}^n y_i}{n} \end{aligned} \quad (1.2)$$

Step 2: Let B be the number of bootstrap samples taken. With $n = 8$ and for each method, sample from the plug-in distribution.

$$\begin{aligned} B_{Y,1}^* &= \{y_{1,1}^*, \dots, y_{1,8}^*\} \stackrel{iid}{\sim} \mathcal{P}(\hat{\lambda}_Y) \\ B_{Y,2}^* &= \{y_{2,1}^*, \dots, y_{2,8}^*\} \stackrel{iid}{\sim} \mathcal{P}(\hat{\lambda}_Y) \\ &\vdots = \vdots \end{aligned} \quad (1.3)$$

$$\begin{aligned} B_{Y,2000}^* &= \{y_{2000,1}^*, \dots, y_{2000,8}^*\} \stackrel{iid}{\sim} \mathcal{P}(\hat{\lambda}_Y) \\ B_{X,1}^* &= \{x_{1,1}^*, \dots, x_{1,8}^*\} \stackrel{iid}{\sim} \mathcal{P}(\hat{\lambda}_X) \\ B_{X,2}^* &= \{x_{2,1}^*, \dots, x_{2,8}^*\} \stackrel{iid}{\sim} \mathcal{P}(\hat{\lambda}_X) \\ &\vdots = \vdots \end{aligned} \quad (1.4)$$

$$B_{X,2000}^* = \{x_{2000,1}^*, \dots, x_{2000,8}^*\} \stackrel{iid}{\sim} \mathcal{P}(\hat{\lambda}_X)$$

Step 3: For each of the bootstrap samples, calculate $\hat{\theta}_b^* = \hat{\mu}_{X,b}^* - \hat{\mu}_{Y,b}^*$; $b = 1, \dots, 2000$.

Step 4: For each of the bootstrap samples, calculate $\hat{\theta}_b^* = \hat{\mu}_{X,b}^* - \hat{\mu}_{Y,b}^*$; $b = 1, \dots, 2000$.


$$\widehat{se}(\hat{\theta}) = \sqrt{\frac{\sum_{i=1}^n (\hat{\theta}_i^* - \bar{\hat{\theta}}^*)^2}{B-1}}; \quad \bar{\hat{\theta}}^* = \frac{1}{B} \sum_{i=1}^B \hat{\theta}_i^* \quad (1.5)$$

```
[1] "The file exists; loading ../../../../problems/midterm/datasets/microbiology.RData"
[1] 7
```

The bootstrap estimate of the standard error is 5.83411.

1.2 Histogram

The histogram has a symmetric bell curve shape. I recalled in class that this is expected to be centered on the statistic. In this case the statistic is 12 while the mean of $\hat{\theta}^*$'s is 12.01794. This is the reason why I chose the mean for the plug-in distribution. If the plug-in distribution is based on the sum, the histogram will be centered on a different, larger value, far from the statistic. It will also have a higher estimate of the standard error (9.92229).



1.3 Confidence Intervals

We are 95% confident that the true value of θ is between 0.625 and 23.375. If instead we use sum for the sufficient statistic, then the CI will be a lot wider, with the sample statistic, not even making it within the interval 73.24688 and 117.625.

2 Fishery

3 References

Sonnet, L. *Standard errors in OLS*. https://lukesonnet.com/teaching/inference/200d_standard_errors.pdf

4 Appendix

4.1 Code to read data for items 1 & 2

```
CODE FILENAME: ../R/01_00_load_data.R
```

```
# wd: /home/scientists/sci01/Projects/bootstrap/solutions/ps_01/child

data <- "../.../problems/midterm/datasets/microbiology.RData"

if (file.exists(data)) {
  print(paste(c("The file exists; loading", data), collapse = ' '))
  load(data)
} else {
  paste(c("The file does not exist; creating, loading and saving", data),
        collapse = ' ')
  microbiology <- data.frame(
    'i' = 1:8,
    'X' = c(176, 125, 152, 180, 159, 168, 160, 151),
    'Y' = c(164, 121, 137, 169, 144, 145, 156, 139)
  )
  n <- dim(microbiology)[1]
  seed <- 7
  B <- 2000
  save(microbiology, n, B, seed, file = data)
}

rm(data)
```